

# Open source database management systems – PostgreSQL & MySQL

Tech blog on PostgreSQL / MySQL for DBAs & Developers

📅 APRIL 18, 2013    👤 PRASHANTH GORIPARTHI

💬 5

## Setup pgbouncer connection pooling for PostgreSQL on CentOS/RedHat/Fedora

This post helps you install, setup and benchmark pgbouncer connection pooling for PostgreSQL 9.2 on CentOS 6/RedHat/Fedora based systems. pgbouncer is one of the PostgreSQL connection poolers. There are other connection poolers available for PostgreSQL. Any enduser applications can be connected to pgbouncer as if it were a PostgreSQL server, and pgbouncer will create a connection to the actual server, or it will reuse one of its existing connections. The purpose of connection poolers in general is to lower the performance impact of opening new connections to PostgreSQL(or Other) databases. pgbouncer can connection pool in three ways, which are

- Session pooling
- Transaction pooling
- Statement pooling

Second and Third options are more aggressive and not very easy to maintain/manage. In this post I will setup session pooling. I am assuming that you have a postgresql database server up and running before you start deploying pgbouncer. If haven't done installing postgresql before, please go through my other post on [how to install Postgres database server](#). To install pgbouncer, simply yum install it as shown below.

```
yum install pgbouncer
```

```
pgoripartn — root@opensourcebms:~ — ssh — 115x34
[root@opensourcebms ~]# yum install pgbouncer
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
 * base: dist1.800hosting.com
 * extras: mirrors.xmission.com
 * updates: mirrors.loosefoot.com
Setting up Install Process
Resolving Dependencies
--> Running transaction check
--> Package pgbouncer.x86_64 0:1.5.4-1.rhel6 will be installed
--> Processing Dependency: libevent-2.0.so.5()(64bit) for package: pgbouncer-1.5.4-1.rhel6.x86_64
--> Running transaction check
--> Package libevent.x86_64 0:2.0.19-1.rhel6 will be installed
--> Finished Dependency Resolution

Dependencies Resolved

=====
Package                                Arch              Version            Repository          Size
=====
Installing:
pgbouncer                             x86_64            1.5.4-1.rhel6      pgdg92              102 k
Installing for dependencies:
libevent                              x86_64            2.0.19-1.rhel6     pgdg92              149 k
=====

Transaction Summary
=====
Install      2 Package(s)

Total download size: 251 k
Installed size: 781 k
Is this ok [y/N]: y
Downloading Packages:
(1/2): libevent-2.0.19-1.rhel6.x86_64.rpm | 149 kB  00:00
(2/2): pgbouncer-1.5.4-1.rhel6.x86_64.rpm | 102 kB  00:00
-----
Total                                     271 kB/s | 251 kB  00:00
Running rpm_check_debug
Running Transaction Test
Transaction Test Succeeded
Running Transaction
  Installing : libevent-2.0.19-1.rhel6.x86_64                1/2
  Installing : pgbouncer-1.5.4-1.rhel6.x86_64                2/2
  Verifying   : libevent-2.0.19-1.rhel6.x86_64                1/2
  Verifying   : pgbouncer-1.5.4-1.rhel6.x86_64                2/2

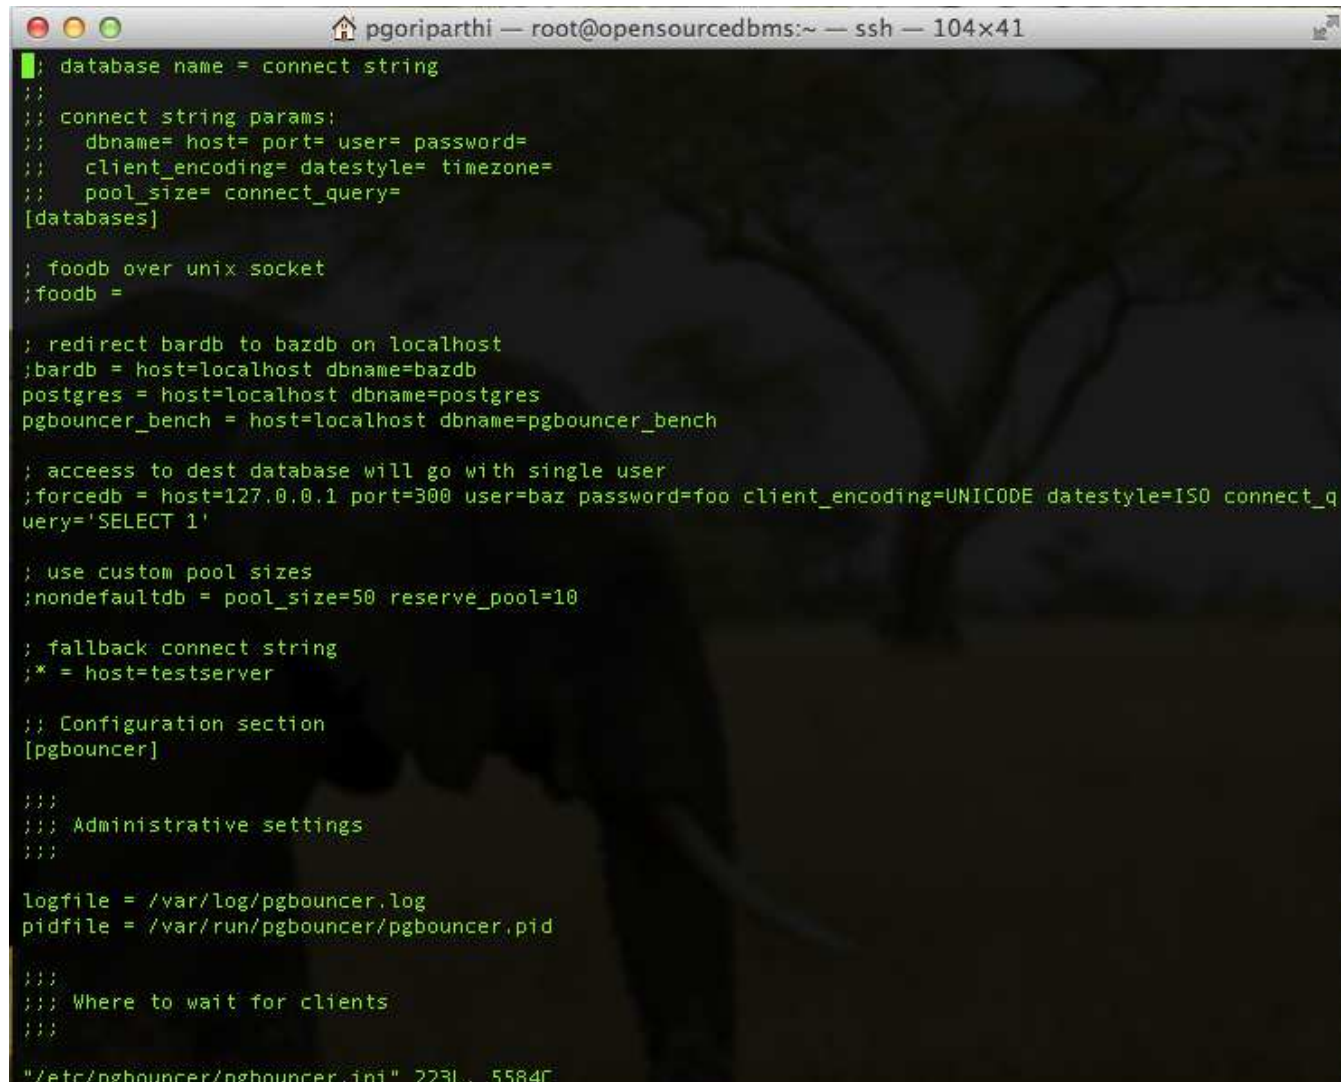
Installed:
pgbouncer.x86_64 0:1.5.4-1.rhel6

Dependency Installed:
libevent.x86_64 0:2.0.19-1.rhel6

Complete!
[root@opensourcebms ~]#
```

Once pgbouncer is installed we need to tweak its config and create some users. Let me walk through these with you. First lets edit the config file and tweak it.

```
vi /etc/pgbouncer/pgbouncer.ini
```

A screenshot of a terminal window with a dark background and light green text. The window title bar shows 'pgoriparthi — root@opensourcebms:~ — ssh — 104x41'. The terminal displays the contents of the file /etc/pgbouncer/pgbouncer.ini. The configuration includes sections for database connection strings, database-specific settings (like foodb, bardb, postgres, pgbouncer\_bench), a forced database connection, pool sizes, a fallback connection string, and administrative settings like logfile and pidfile.

```
#!/usr/bin/perl
; database name = connect string
;;
;; connect string params:
;;   dbname= host= port= user= password=
;;   client_encoding= datestyle= timezone=
;;   pool_size= connect_query=
[databases]

; foodb over unix socket
;foodb =

; redirect bardb to bazdb on localhost
;bardb = host=localhost dbname=bazdb
postgres = host=localhost dbname=postgres
pgbouncer_bench = host=localhost dbname=pgbouncer_bench

; access to dest database will go with single user
;forcedb = host=127.0.0.1 port=300 user=baz password=foo client_encoding=UNICODE datestyle=ISO connect_q
uery='SELECT 1'

; use custom pool sizes
;nondefaultdb = pool_size=50 reserve_pool=10

; fallback connect string
;* = host=testserver

;; Configuration section
[pgbouncer]

;;;
;;; Administrative settings
;;;

logfile = /var/log/pgbouncer.log
pidfile = /var/run/pgbouncer/pgbouncer.pid

;;;
;;; Where to wait for clients
;;;

"/etc/pgbouncer/pgbouncer.ini" 223L, 5584C
```

You will need to update/add the following lines to your config:

Add all the databases that you need to pass through from pgbouncer, I just added postgres/pgbouncer\_bench db redirects here (I will be using pgbouncer\_bench database down for benchmarking). Also you can change the auth\_type to your needs (I am using plain for illustration purpose).

```
postgres = host=localhost dbname=postgres
pgbouncer_bench = host=localhost dbname=pgbouncer_bench

listen_addr = *
auth_file = /etc/pgbouncer/userlist.txt
auth_type = plain
```

Your full config file should look as following:

```
;; database name = connect string
;;
;; connect string params:
;;   dbname= host= port= user= password=
;;   client_encoding= datestyle= timezone=
;;   pool_size= connect_query=
[databases]

; foodb over unix socket
;foodb =

; redirect bardb to bazdb on localhost
;bardb = host=localhost dbname=bazdb
postgres = host=localhost dbname=postgres
pgbouncer_bench = host=localhost dbname=pgbouncer_bench

; access to dest database will go with single user
;forcedb = host=127.0.0.1 port=300 user=baz password=foo client_encoding=UNICODE datestyle=ISO connect_query='SELECT'

; use custom pool sizes
;nondefaultdb = pool_size=50 reserve_pool=10
```

```
; fallback connect string
;* = host=testserver

;; Configuration section
[pgbouncer]

;;;
;;; Administrative settings
;;;

logfile = /var/log/pgbouncer.log
pidfile = /var/run/pgbouncer/pgbouncer.pid

;;;
;;; Where to wait for clients
;;;

; ip address or * which means all ip-s
listen_addr = *
listen_port = 6432

; unix socket is also used for -R.
; On debian it should be /var/run/postgresql
;unix_socket_dir = /tmp
;unix_socket_mode = 0777
;unix_socket_group =

;;;
;;; Authentication settings
;;;

; any, trust, plain, crypt, md5
auth_type = plain
;auth_file = /8.0/main/global/pg_auth
auth_file = /etc/pgbouncer/userlist.txt

;;;
;;; Users allowed into database 'pgbouncer'
;;;

; comma-separated list of users, who are allowed to change settings
admin_users = postgres

; comma-separated list of users who are just allowed to use SHOW command
stats_users = stats, postgres
```

```
;;;
;;; Pooler personality questions
;;;

; When server connection is released back to pool:
;   session      - after client disconnects
;   transaction  - after transaction finishes
;   statement    - after statement finishes
pool_mode = session

;
; Query for cleaning connection immediately after releasing from client.
; No need to put ROLLBACK here, pgbouncer does not reuse connections
; where transaction is left open.
;
; Query for 8.3+:
;   DISCARD ALL;
;
; Older versions:
;   RESET ALL; SET SESSION AUTHORIZATION DEFAULT
;
; Empty if transaction pooling is in use.
;
server_reset_query = DISCARD ALL

;
; Comma-separated list of parameters to ignore when given
; in startup packet.  Newer JDBC versions require the
; extra_float_digits here.
;
;ignore_startup_parameters = extra_float_digits

;
; When taking idle server into use, this query is ran first.
;   SELECT 1
;
;server_check_query = select 1

; If server was used more recently that this many seconds ago,
; skip the check query.  Value 0 may or may not run in immediately.
;server_check_delay = 30

;;;
;;; Connection limits
;;;
```

```
; total number of clients that can connect
max_client_conn = 100

; default pool size. 20 is good number when transaction pooling
; is in use, in session pooling it needs to be the number of
; max clients you want to handle at any moment
default_pool_size = 100

; how many additional connection to allow in case of trouble
;reserve_pool_size = 5

; if a clients needs to wait more than this many seconds, use reserve pool
;reserve_pool_timeout = 3

; log if client connects or server connection is made
;log_connections = 1

; log if and why connection was closed
;log_disconnections = 1

; log error messages pooler sends to clients
;log_pooler_errors = 1

; If off, then server connections are reused in LIFO manner
;server_round_robin = 0

;;;
;;; Timeouts
;;;

;; Close server connection if its been connected longer.
;server_lifetime = 1200

;; Close server connection if its not been used in this time.
;; Allows to clean unnecessary connections from pool after peak.
;server_idle_timeout = 60

;; Cancel connection attempt if server does not answer takes longer.
;server_connect_timeout = 15

;; If server login failed (server_connect_timeout or auth failure)
;; then wait this many second.
;server_login_retry = 15

;; Dangerous. Server connection is closed if query does not return
;; in this time. Should be used to survive network problems,
```

```
;; _not_ as statement_timeout. (default: 0)
;query_timeout = 0

;; Dangerous. Client connection is closed if the query is not assigned
;; to a server in this time. Should be used to limit the number of queued
;; queries in case of a database or network failure. (default: 0)
;query_wait_timeout = 0

;; Dangerous. Client connection is closed if no activity in this time.
;; Should be used to survive network problems. (default: 0)
;client_idle_timeout = 0

;; Disconnect clients who have not managed to log in after connecting
;; in this many seconds.
;client_login_timeout = 60

;; Clean automatically created database entries (via "") if they
;; stay unused in this many seconds.
;autodb_idle_timeout = 3600

;;;
;;; Low-level tuning options
;;;

;; buffer for streaming packets
;pkt_buf = 2048

;; man 2 listen
;listen_backlog = 128

;; networking options, for info: man 7 tcp

;; Linux: notify program about new connection only if there
;; is also data received. (Seconds to wait.)
;; On Linux the default is 45, on other OS'es 0.
;tcp_defer_accept = 0

;; In-kernel buffer size (Linux default: 4096)
;tcp_socket_buffer = 0

;; whether tcp keepalive should be turned on (0/1)
;tcp_keepalive = 1

;; following options are Linux-specific.
;; they also require tcp_keepalive=1
```



```
;; count of keepaliva packets
;tcp_keepcnt = 0

;; how long the connection can be idle,
;; before sending keepalive packets
;tcp_keepidle = 0

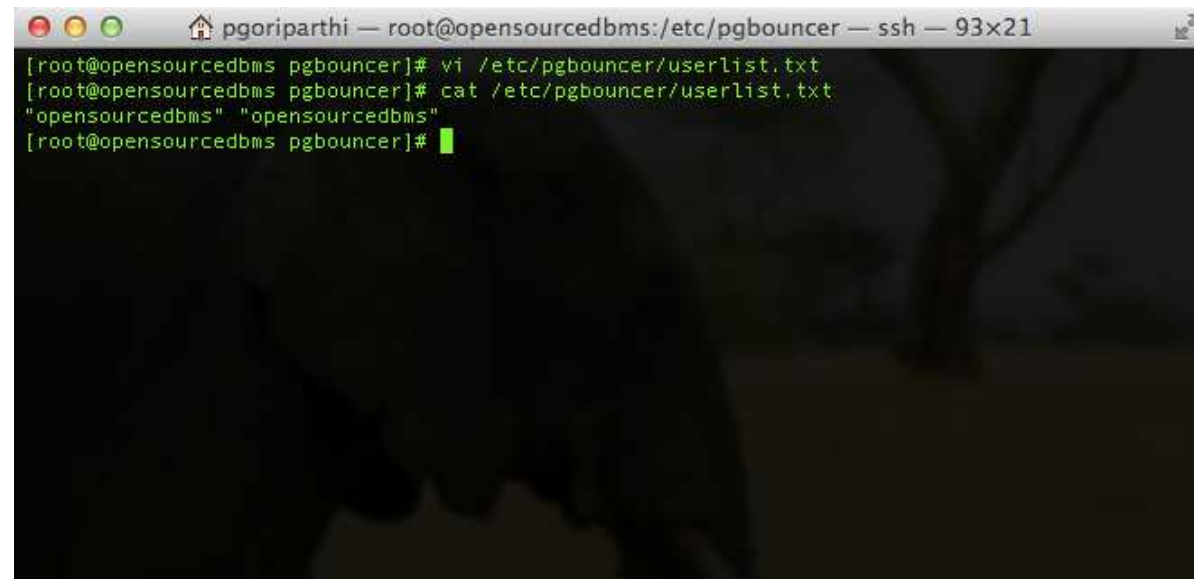
;; The time between individual keepalive probes.
;tcp_keepintvl = 0

;; DNS lookup caching time
;dns_max_ttl = 15

;; DNS zone SOA lookup period
;dns_zone_check_period = 0
```

Now lets create users.txt file for accessing database through pgbouncer. These users are users that already exist in our existing database. Add all users you wish to give access to your databases through pgbouncer. Format of user/password should be like "user" "password".

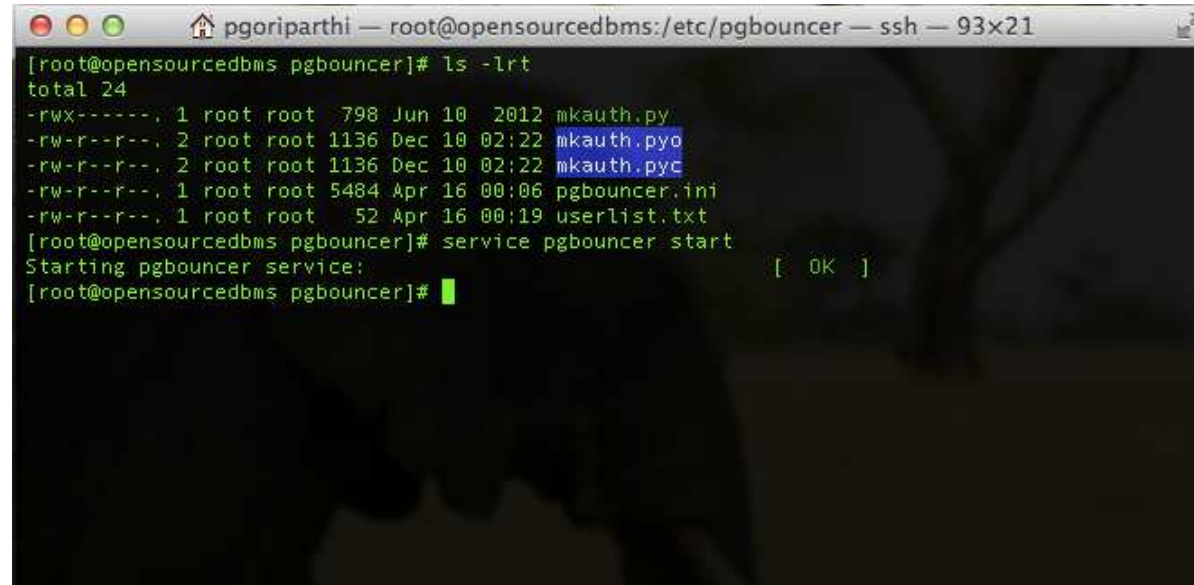
```
vi /etc/pgbouncer/userlist.txt
```

A terminal window titled 'pgoriparthi — root@opensourcedbms:/etc/pgbouncer — ssh — 93x21'. The terminal shows the following commands and output:

```
[root@opensourcedbms pgbouncer]# vi /etc/pgbouncer/userlist.txt
[root@opensourcedbms pgbouncer]# cat /etc/pgbouncer/userlist.txt
"opensourcedbms" "opensourcedbms"
[root@opensourcedbms pgbouncer]#
```

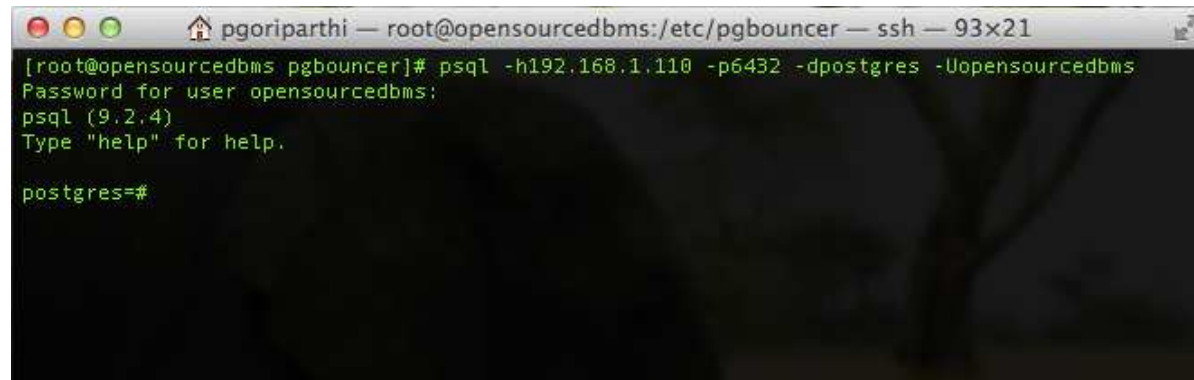
Now that we created users we can start our connection pooler.

```
service pgbouncer start
```

A terminal window titled 'pgoriparthi — root@opensourcebms:/etc/pgbouncer — ssh — 93x21'. The user runs 'ls -lrt' in the '[root@opensourcebms pgbouncer]# ' prompt, showing a list of files: 'mkauth.py', 'mkauth.pyo', 'mkauth.pyc', 'pgbouncer.ini', and 'userlist.txt'. Then, the user runs 'service pgbouncer start', which outputs 'Starting pgbouncer service: [ OK ]'. The prompt returns to '[root@opensourcebms pgbouncer]# ' with a green cursor.


```
[root@opensourcebms pgbouncer]# ls -lrt
total 24
-rwx-----, 1 root root  798 Jun 10  2012 mkauth.py
-rw-r--r--, 2 root root 1136 Dec 10  02:22 mkauth.pyo
-rw-r--r--, 2 root root 1136 Dec 10  02:22 mkauth.pyc
-rw-r--r--, 1 root root 5484 Apr 16  00:06 pgbouncer.ini
-rw-r--r--, 1 root root   52 Apr 16  00:19 userlist.txt
[root@opensourcebms pgbouncer]# service pgbouncer start
Starting pgbouncer service: [ OK ]
[root@opensourcebms pgbouncer]#
```

Thats it folks you have pgbouncer up and running. Now lets connect to it at port 6432 [default, this can be updated in pgbouncer.ini file].

A terminal window titled 'pgoriparthi — root@opensourcebms:/etc/pgbouncer — ssh — 93x21'. The user runs 'psql -h192.168.1.110 -p6432 -dpostgres -Uopensourcebms'. The prompt changes to 'postgres=#'.

```
[root@opensourcebms pgbouncer]# psql -h192.168.1.110 -p6432 -dpostgres -Uopensourcebms
Password for user opensourcedbms:
psql (9.2.4)
Type "help" for help.

postgres=#
```



You have now successfully connected to pgbouncer. Lets benchmark and see if pgbouncer makes any difference or not. We need pgbench utility to benchmark postgres. To get that utility we need to install postgres contrib package.

```
yum install postgresql92-contrib
```

```
[root@opensourcedbms ~]# yum install postgresql92-contrib
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
 * base: dist1.800hosting.com
 * extras: mirrors.xmission.com
 * updates: mirrors.loosefoot.com
Setting up Install Process
Resolving Dependencies
--> Running transaction check
--> Package postgresql92-contrib.x86_64 0:9.2.4-1PGDG.rhel6 will be installed
--> Finished Dependency Resolution

Dependencies Resolved

=====
Package                        Arch          Version           Repository        Size
=====
Installing:
 postgresql92-contrib          x86_64        9.2.4-1PGDG.rhel6 pgdg92            435 k
=====

Transaction Summary
=====
Install      1 Package(s)

Total download size: 435 k
Installed size: 1.6 M
Is this ok [y/N]: y
Downloading Packages:
 postgresql92-contrib-9.2.4-1PGDG.rhel6.x86_64.rpm | 435 kB    00:00
Running rpm_check_debug
Running Transaction Test
Transaction Test Succeeded
Running Transaction
```

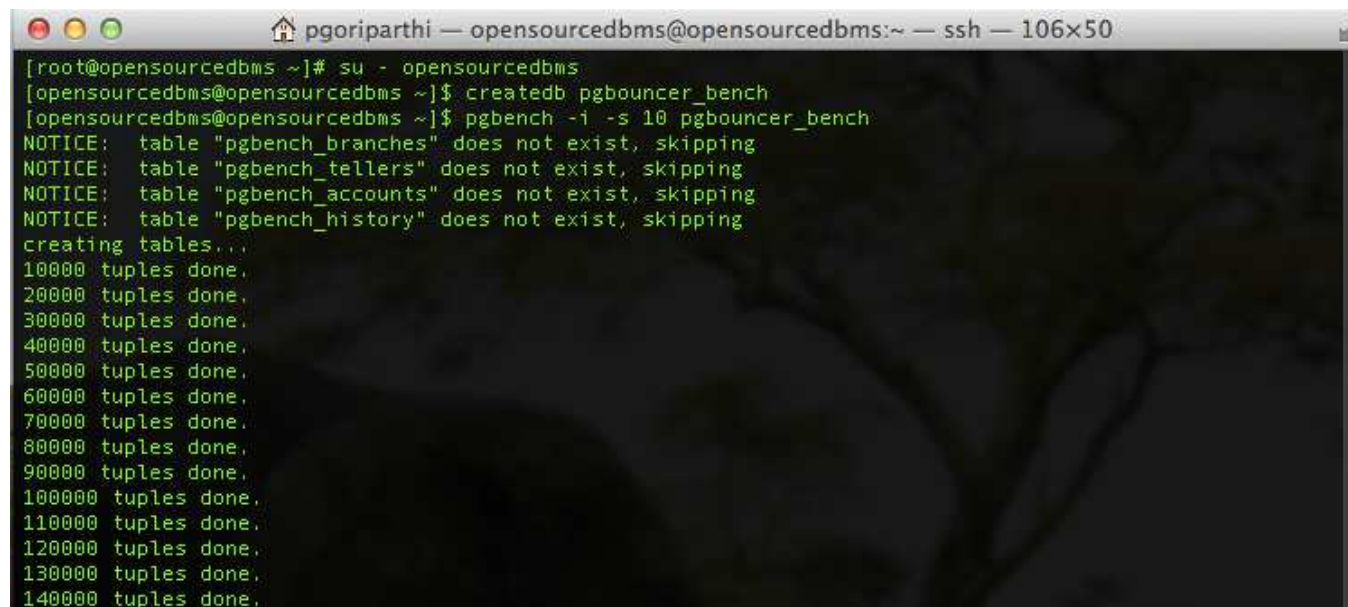
```
Installing : postgresql92-contrib-9.2.4-1PGDG.rhel6.x86_64 1/1
Verifying : postgresql92-contrib-9.2.4-1PGDG.rhel6.x86_64 1/1

Installed:
  postgresql92-contrib.x86_64 0:9.2.4-1PGDG.rhel6

Complete!
[root@opendbms ~]#
```

To benchmark we need to create pgbouncer\_bench(temporary) database to run our connections against and populate this database with some data.

```
createdb pgbouncer_bench
pgbench -i -s 10 pgbouncer_bench
```

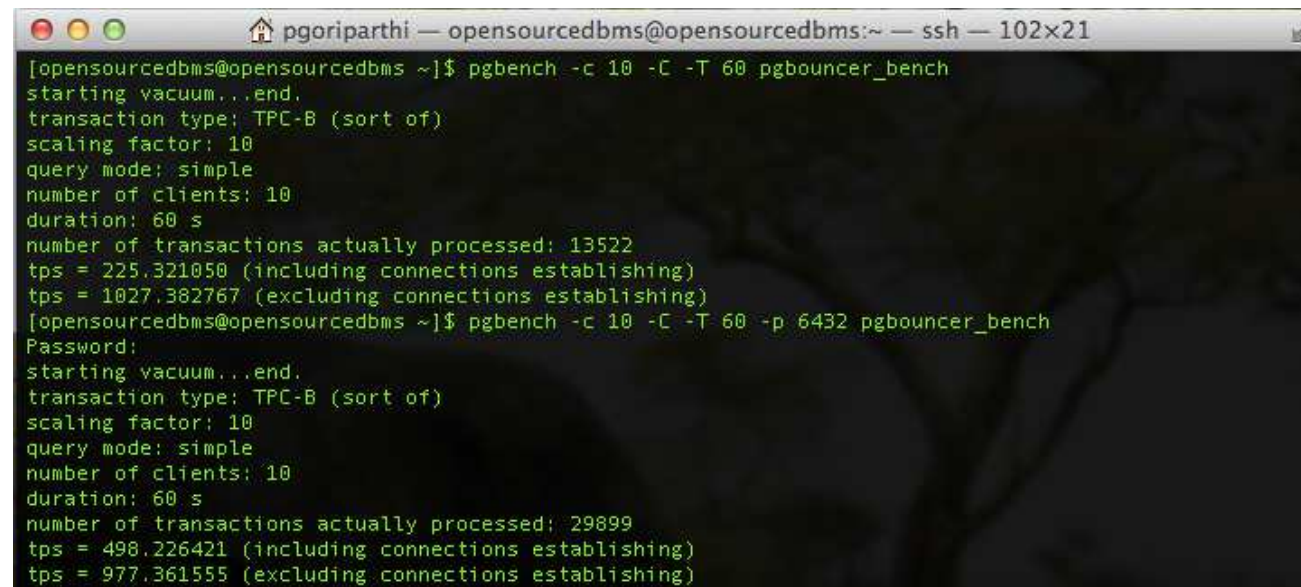


```
pgoriparthi — opendbms@opendbms:~ — ssh — 106x50
[root@opendbms ~]# su - opendbms
[opendbms@opendbms ~]$ createdb pgbouncer_bench
[opendbms@opendbms ~]$ pgbench -i -s 10 pgbouncer_bench
NOTICE: table "pgbench_branches" does not exist, skipping
NOTICE: table "pgbench_tellers" does not exist, skipping
NOTICE: table "pgbench_accounts" does not exist, skipping
NOTICE: table "pgbench_history" does not exist, skipping
creating tables...
10000 tuples done.
20000 tuples done.
30000 tuples done.
40000 tuples done.
50000 tuples done.
60000 tuples done.
70000 tuples done.
80000 tuples done.
90000 tuples done.
100000 tuples done.
110000 tuples done.
120000 tuples done.
130000 tuples done.
140000 tuples done.
```

Now we can start benchmarking our database between port 5432 and 6432(pgbouncer) and see the performance difference.

Performance at 10 client connections:

```
[opentosourcedbms@opentosourcedbms ~]$ pgbench -c 10 -C -T 60 pgbouncer_bench
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 10
duration: 60 s
number of transactions actually processed: 13522
tps = 225.321050 (including connections establishing)
tps = 1027.382767 (excluding connections establishing)
[opentosourcedbms@opentosourcedbms ~]$ pgbench -c 10 -C -T 60 -p 6432 pgbouncer_bench
Password:
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 10
duration: 60 s
number of transactions actually processed: 29899
tps = 498.226421 (including connections establishing)
tps = 977.361555 (excluding connections establishing)
```

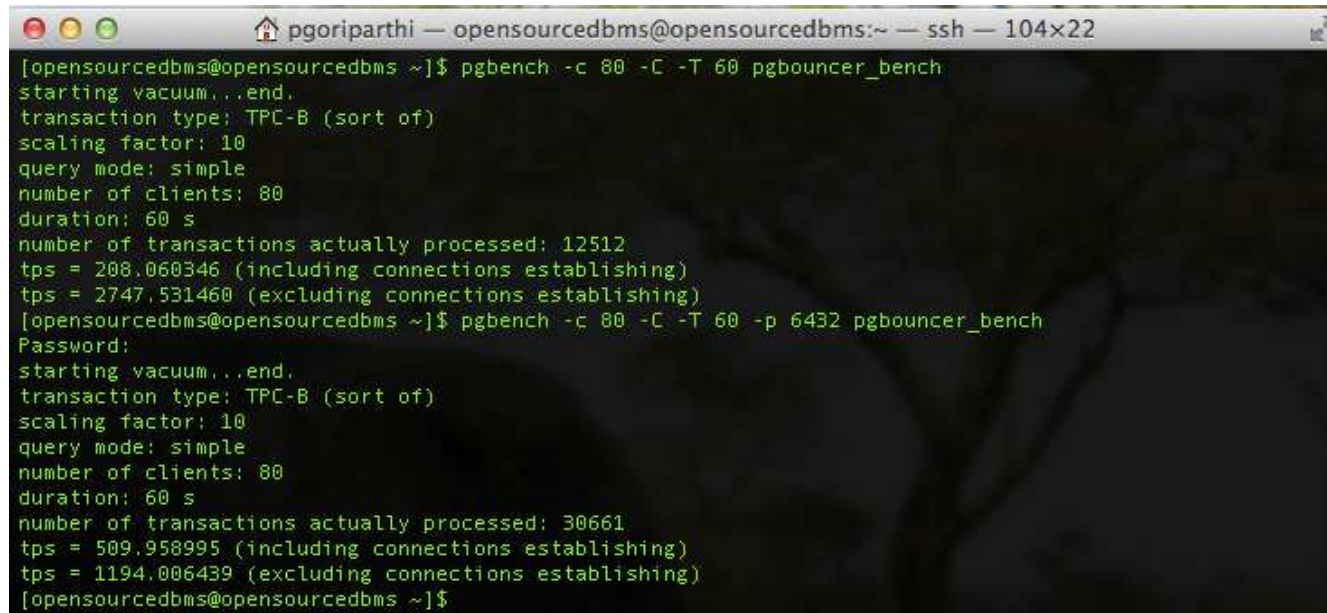
A screenshot of a terminal window titled "pgoriparthi — opentosourcedbms@opentosourcedbms:~ — ssh — 102x21". The terminal displays the output of two pgbench commands. The first command is "pgbench -c 10 -C -T 60 pgbouncer\_bench", which shows a transaction type of TPC-B (sort of), a scaling factor of 10, and a duration of 60 seconds. It reports that 13522 transactions were actually processed, with a tps of 225.321050 (including connections establishing) and 1027.382767 (excluding connections establishing). The second command is "pgbench -c 10 -C -T 60 -p 6432 pgbouncer\_bench", which prompts for a password and then shows the same transaction type and scaling factor. It reports that 29899 transactions were actually processed, with a tps of 498.226421 (including connections establishing) and 977.361555 (excluding connections establishing).

```
pgoriparthi — opentosourcedbms@opentosourcedbms:~ — ssh — 102x21
[opentosourcedbms@opentosourcedbms ~]$ pgbench -c 10 -C -T 60 pgbouncer_bench
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 10
duration: 60 s
number of transactions actually processed: 13522
tps = 225.321050 (including connections establishing)
tps = 1027.382767 (excluding connections establishing)
[opentosourcedbms@opentosourcedbms ~]$ pgbench -c 10 -C -T 60 -p 6432 pgbouncer_bench
Password:
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 10
duration: 60 s
number of transactions actually processed: 29899
tps = 498.226421 (including connections establishing)
tps = 977.361555 (excluding connections establishing)
```

Performance at 80 client connections:



```
[opensourcedbms@opensourcedbms ~]$ pgbench -c 80 -C -T 60 pgbouncer_bench
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 80
duration: 60 s
number of transactions actually processed: 12512
tps = 208.060346 (including connections establishing)
tps = 2747.531460 (excluding connections establishing)
[opensourcedbms@opensourcedbms ~]$ pgbench -c 80 -C -T 60 -p 6432 pgbouncer_bench
Password:
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 80
duration: 60 s
number of transactions actually processed: 30661
tps = 509.958995 (including connections establishing)
tps = 1194.006439 (excluding connections establishing)
```



```
pgoriparthi — opensourcedbms@opensourcedbms:~ — ssh — 104x22
[opensourcedbms@opensourcedbms ~]$ pgbench -c 80 -C -T 60 pgbouncer_bench
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 80
duration: 60 s
number of transactions actually processed: 12512
tps = 208.060346 (including connections establishing)
tps = 2747.531460 (excluding connections establishing)
[opensourcedbms@opensourcedbms ~]$ pgbench -c 80 -C -T 60 -p 6432 pgbouncer_bench
Password:
starting vacuum...end.
transaction type: TPC-B (sort of)
scaling factor: 10
query mode: simple
number of clients: 80
duration: 60 s
number of transactions actually processed: 30661
tps = 509.958995 (including connections establishing)
tps = 1194.006439 (excluding connections establishing)
[opensourcedbms@opensourcedbms ~]$
```

You can see the significant improvement of transactions when we are connecting to database using pgbouncer. There are various tweaks that you can do to pgbouncer config file and tune according to your needs.

Please post questions in comments if you have any.



📁 POSTGRESQL

## About Prashanth Goriparthi

Sr. Database Administrator / Architect [View all posts by Prashanth Goriparthi →](#)

## 5 Comments

**Tiah**

August 8, 2013 at 9:40 am

Nice article, very useful 😊

↪ Reply

**Shariman**

March 19, 2014 at 6:13 am

Very detail and accurate guide. Thanks a lot

↪ Reply

[inetholic](#)

March 31, 2014 at 5:28 pm

Thanks ... i can increase TPS with pgbouncer ... good job bro ...

➔ Reply

**Akbar**

February 4, 2015 at 3:50 am

Hai. I have some error.

I was successfully run this :

```
[postgres@edb ~]$ pgbench -c 10 -C -T 60 pgbouncer_bench
```

But when I run this :

```
[postgres@edb ~]$ pgbench -c 10 -C -T 60 -p 6432 pgbouncer_bench
```

Getting error :

Connection to database "pgbouncer\_bench" failed:

ERROR: No such database: pgbouncer\_bench

Do you have idea why this is happen?

Thanks

➔ Reply

**Jagdeep**

December 22, 2015 at 4:19 pm

I am facing below mentioned error after following the above mentioned steps.

psql: ERROR: pgbouncer cannot connect to server

Could you please able to help me to resolve this issue?

Thanks in advance.

➔ Reply



## Leave a Reply


Name:

Email:

Website

\* Captcha [Please fill/solve below] : \*

Time limit is exhausted. Please reload the CAPTCHA.

5 ×  = twenty 

Submit Comment

Open source database management systems - PostgreSQL & MySQL © 2017

[Privacy Policy](#) - [Terms and Conditions](#)